

Introduction to R

Greg Ridgeway

2026-07-13

Table of contents

1	Introduction	2
2	Basic Math and Functions in R	2
3	Combining values together into a collection (or vector)	3
4	Things R will complain about	7
5	Setting the working directory	9
6	Loading a first dataset, protests in the United States	10
7	Indexing basics	15
8	Creating objects	194
8.1	Exercises	195
9	Filtering rows	195
9.1	Exercise	198
10	Summarizing data	198
11	Mutate to edit and create new columns	200
11.1	Exercises	208
12	Creating your own functions	208
13	Summarizing with groups of protests	211
13.1	Exercises	215
14	Pivot wider and longer	216

15 Graphics and plots	222
16 Review	232
17 Solutions to the exercises	233

1 Introduction

This is the first set of notes for an introduction to R programming from criminology and criminal justice. These notes assume that you have the latest version of R and RStudio installed. We are also assuming that you know how to start a new script file and submit code to the R console. From that basic knowledge about using R, we are going to start with `2+2` and by the end of this set of notes you will load in a dataset on protests in the United States (mostly), create a few plots, count some incidents, and be able to do some basic data manipulations. Our aim is to build a firm foundation on which we will build throughout this set of notes.

R sometimes provides useful help as to how to do something, such as choosing the right function or figuring what the syntax of a line of code should be. Let's say we're stumped as to what the `sqrt()` function does. Just type `?sqrt` at the R prompt to read documentation on `sqrt()`. Most help pages have examples at the bottom that can give you a better idea about how the function works. R has over 7,000 functions and an often seemingly inconsistent syntax. As you do more complex work with R (such as using new packages), the Help tab can be useful.

2 Basic Math and Functions in R

R, on a very unsophisticated level, is like a calculator.

```
2+2
1*2*3*4
(1+2+3-4)/(5*7)
sqrt(2)
(1+sqrt(5))/2 # golden ratio
2^3
log(2.718281828)
round(2.718281828,3)
12^2
factorial(4)
abs(-4)
```

```
[1] 4
```

```
[1] 24
[1] 0.05714286
[1] 1.414214
[1] 1.618034
[1] 8
[1] 1
[1] 2.718
[1] 144
[1] 24
[1] 4
```

3 Combining values together into a collection (or vector)

We will use the `c()` function a lot. `c()` combines elements, like numbers and text to form a vector or a collection of values. If we wanted to combine the numbers 1 to 5 we could do

```
c(1,2,3,4,5)
```

```
[1] 1 2 3 4 5
```

With the `c()` function, it's important to separate all of the items with commas.

Conveniently, if you want to add 1 to each item in this collection, there's no need to add 1 like `c(1+1,2+1,3+1,4+1,5+1)`... that's a lot of typing. Instead R offers the shortcut

```
c(1,2,3,4,5)+1
```

```
[1] 2 3 4 5 6
```

In fact, you can apply any mathematical operation to each value in the same way.

```
c(1,2,3,4,5)*2
sqrt(c(1,2,3,4,5))
(c(1,2,3,4,5)-3)^2
abs(c(-1,1,-2,2,-3,3))
```

```
[1] 2 4 6 8 10
[1] 1.000000 1.414214 1.732051 2.000000 2.236068
[1] 4 1 0 1 4
[1] 1 1 2 2 3 3
```

Note in the examples below that you can also have a collection of non-numerical items. When combining text items, remember to use quotes around each item.

```
c("CRIM6000", "CRIM6001", "CRIM6002", "CRIM6003")  
c("yes", "no", "no", NA, NA, "yes")
```

```
[1] "CRIM6000" "CRIM6001" "CRIM6002" "CRIM6003"  
[1] "yes" "no" "no" NA NA "yes"
```

In R, NA means a missing value. We will do more exercises later using data containing some NA values. In any dataset in the wild, you are virtually guaranteed to find some NAs. The function `is.na()` helps determine whether there are any missing values (any NAs). In some of the problems below, we will use `is.na()`.

One strange but important thing about NA is that R treats it as “unknown.” If you ask whether one unknown value equals another unknown value, R still cannot know the answer. This surprises a lot of new R users.

```
NA == NA
```

```
[1] NA
```

```
is.na(NA)
```

```
[1] TRUE
```

Use `is.na()` when you want to test whether a value is missing. Do not use `== NA`.

You can use double quotes or single quotes in R as long as you are consistent. When you have quotes inside the text you need to be particularly careful.

```
"Lou Gehrig's disease"  
'The officer shouted "halt!"'
```

```
[1] "Lou Gehrig's disease"  
[1] "The officer shouted \"halt!\""
```

The backslashes in the above text “protect” the double quote, communicating to you and to R that the next double quote is not the end of the text, but a character that is actually part of the text you want to keep.

The `c()` function is not the only way to make a collection of values in R. For example, placing a `:` between two numbers can return a collection of numbers in sequence. The functions `rep()` and `seq()` produce repeated values or sequences.

```
1:10
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
5:-5
```

```
[1] 5 4 3 2 1 0 -1 -2 -3 -4 -5
```

```
c(1,1,1,1,1,1,1,1,1,1)
```

```
[1] 1 1 1 1 1 1 1 1 1 1
```

```
rep(1,10)
```

```
[1] 1 1 1 1 1 1 1 1 1 1
```

```
rep(c(1,2),each=5)
```

```
[1] 1 1 1 1 1 2 2 2 2 2
```

```
seq(1, 5)
```

```
[1] 1 2 3 4 5
```

```
seq(1, 5, 2)
```

```
[1] 1 3 5
```

R will also do arithmetic with two vectors, doing the calculation pairwise. The following will compute $1+11$ and $2+12$ up to $10+20$.

```
1:10 + 11:20
```

```
[1] 12 14 16 18 20 22 24 26 28 30
```

Yet, other functions operate on the whole collection of values in a vector. See the following examples:

```
sum(c(1,10,3,6,2,5,8,4,7,9)) # sum
```

```
[1] 55
```

```
length(c(1,10,3,6,2,5,8,4,7,9)) # how many?
```

```
[1] 10
```

```
cumsum(c(1,10,3,6,2,5,8,4,7,9)) # cumulative sum
```

```
[1] 1 11 14 20 22 27 35 39 46 55
```

```
mean(c(1,10,3,6,2,5,8,4,7,9)) # mean of collection of 10 numbers
```

```
[1] 5.5
```

```
median(c(1,10,3,6,2,5,8,4,7,9)) # median of same population
```

```
[1] 5.5
```

There are also some functions in R that help us find the biggest and smallest values. For example:

```
max(c(1,10,3,6,2,5,8,4,7,9)) # what is the biggest value in vector?
```

```
[1] 10
```

```
which.max(c(1,10,3,6,2,5,8,4,7,9)) # in which "spot" would we find it?
```

```
[1] 2
```

```
min(c(1,10,3,6,2,5,8,4,7,9)) # what is the smallest value in vector?
```

```
[1] 1
```

```
which.min(c(1,10,3,6,2,5,8,4,7,9)) # in which "spot" would we find it?
```

```
[1] 1
```

4 Things R will complain about

When you are learning R, you will see error messages. This is normal. Error messages are R's way of saying, "I cannot do what you asked me to do yet."

R is case sensitive. That means **Attendees**, **attendees**, and **ATTENDEES** are three different names. If you ask R for an object that does not exist, it will complain that the object is not found.

```
attendees
```

Error:

```
! object 'attendees' not found
```

R also needs all function inputs to be separated by commas. If a comma is missing, R often says something like "unexpected numeric constant" or "unexpected symbol."

```
c(1 2 3)
```

```
Error in parse(text = input): <text>:1:5: unexpected numeric constant
```

```
1: c(1 2
    ^
```

R needs opening and closing parentheses to match. If you start a function with `(`, you need to close it with `)`. If you neglect the closing parenthesis, like `sqrt(25`, R cannot finish reading the code. In the console, it will usually display a `+`, R's way of saying "I'm waiting for you to finish your thought. You're leaving me hanging here." When rendering this document, R reports it as an incomplete expression.

```
> sqrt(25
+
```

To get out of this, either type the missing `)` or hit escape to abandon the calculation.

Text values need quotes. Without quotes, R thinks you are naming an object.

```
c(yes, no, yes)
```

```
Error:
! object 'yes' not found
```

File names and folder names have to be exactly right. If R says that it cannot open a file or that a file does not exist, check the spelling, the folder, and the working directory. (We will load this file later, but here it is missing an "s" in `protests.RData`)

```
load("protest.RData")
```

```
Warning in readChar(con, 5L, useBytes = TRUE): cannot open compressed file
'protest.RData', probable reason 'No such file or directory'
```

```
Error in `readChar()`:
! cannot open the connection
```

The fastest debugging strategy is usually to read the error message, check spelling and capitalization, check commas and parentheses, and then run just one line at a time.

5 Setting the working directory

Now that we have covered a lot of fundamental R features, it is time to load in a real dataset. However, before we do that, R needs to know where to find the data file. So we first need to talk about “the working directory”. When you start R, it has a default folder or directory on your computer where it will retrieve or save any files. You can run `getwd()` to get the current working directory. Here’s my current working directory, which will not be the same as yours.

```
getwd()
```

```
[1] "C:/gitdev/R4crim"
```

Almost certainly this default directory is *not* where you plan to have all of your datasets and files stored. Instead, you probably have an “analysis” or “project” or “R4crim” folder somewhere on your computer where you would like to store your data and work.

Use `setwd()` to tell R what folder you want it to use as the working directory. If you do not set the working directory, R will not know where to find the data you wish to import and will save your results in a location in which you would probably never look. For this example, we are going to use `setwd()` as the first line of this script. Later we will talk about setting up a more portable R Studio project.

```
setwd("C:/Users/greg_/CRIM6002/notes/R4crim")
```

Note that for all platforms, Windows, Macs, and Linux, the working directory only uses forward slashes. So Windows users be careful... most Windows applications use backslashes, but in an effort to make R scripts work across all platforms, R requires forward slashes. Backslashes have a different use in R that you will meet later.

If you do not know how to write your working directory, here comes RStudio to the rescue. In RStudio click Session -> Set Working Directory -> Choose Directory. Then click through to navigate to the working directory that you want to use. When you find it click “Select Folder”. Then look over at the console. RStudio will construct the right `setwd()` syntax for you. Copy and paste that into your script for use later. No need to have to click through the Session menu again now that you have your `setwd()` set up.

Now you can use R functions to load in any datasets that are in your working folder. If you have done your `setwd()` correctly, you should not get any errors because R will know exactly where to look for the data files. If the working directory that you have given in the `setwd()` is not right, R will think the file does not even exist. For example, if you give the path for, say, your R4econ folder, R will not be able to load data because the file is not stored in what R thinks is your working directory. With that out of the way, let’s load a dataset.

6 Loading a first dataset, protests in the United States

We are going to use a dataset of protests in the United States. The data comes from [CountLove](#). The data is a collection of protests that occurred in the United States from 2017 through January 2021. The data includes the date of the protest, the location, the number of attendees, and the reason for the protest. We will load the data and explore it. They stopped collection in February 2021, but you can find more recent crowd data at the [Crowd Counting Consortium](#).

We start by loading in the dataset. I have created a `.RData` file containing the protest data. This is stored in a special format that R can read quickly. The file is called `protests.RData`. We will load this file into R using the `load()` function. Once we have loaded the data, we can see what is in the dataset using the `ls()` function. This will list all the objects in the current environment. If you have just started using R, most likely the only object you see in your environment is `dataProtest`.

```
load("protests.RData")
ls()
```

```
[1] "dataProtest"
```

To start exploring the protest data, have a look at how many rows (protests) and how many columns (protest features) are in the dataset. Then use the `head()` function to show the first few rows of the dataset.

```
# how many rows?
nrow(dataProtest)
```

```
[1] 38097
```

```
# how many columns?
ncol(dataProtest)
```

```
[1] 8
```

```
head(dataProtest)
```

	Date	Location	Attendees
1	2017-01-15	Bowie State University, Bowie, MD	1500
2	2017-01-16	Johnson City, TN	300
3	2017-01-16	Indianapolis, IN	20

4	2017-01-16	Cincinnati, OH	NA
5	2017-01-18	Hartford, CT	300
6	2017-01-19	Washington, DC	NA

Event..legacy..see.tags.

1	Healthcare
2	Civil Rights
3	Environment
4	Other (Martin Luther King Jr.)
5	Healthcare (Pro-Planned Parenthood)
6	Executive

Tags Curated

1	Healthcare; For Affordable Care Act	Yes
2	Civil Rights; For racial justice; Martin Luther King, Jr.	Yes
3	Environment; For wilderness preservation	Yes
4	Civil Rights; For racial justice; Martin Luther King, Jr.	Yes
5	Healthcare; For Planned Parenthood	Yes
6	Executive; Against 45th president	Yes

1	http://www.capitalgazette.com/news-local/ac-cn-aca-rally-0116-20170115-story.html
2	http://www.johnsoncitypress.com/Local/2017/01/16/Hundreds-march-to-honor-Martin-Luther-King-Jr-in-Johnson-City.html?ci=featured&lp=46&ti=
3	http://wishtv.com/2017/01/16/nature-groups-protest-deforestation-of-crown-hill-cemetery-land/
4	http://www.cincinnati.com/picture-gallery/news/2017/01/16/mlk-coalition-march-downtown/96636250/
5	http://www.realhartford.org/freedom-day-of-action/
6	https://malvern-online.com/content/melee-near-trumps-parade-route-police-protesters-clash

Total.Articles

1	1
2	4
3	1
4	1
5	1
6	1

We learn that the dataset has 38097 rows and 8 columns. The `head()` function shows the first few rows of the dataset. The first column is the date of the protest (**Date**), the second is the location (**Location**), and the third is the number of attendees (**Attendees**). The fifth column contains tags describing the purpose of the protest (**Tags**). The other columns contain

other details, like links to news articles about the protest. We will not be using these other features.

Some R functionality relies on packages written by others. For certain basic data tasks, such as selecting certain columns, filtering rows, modifying values, and summarizing data, we will use the `dplyr` package (usually pronounced dee-ply-er... intended to evoke pliers for data). If you do not have `dplyr` installed, you can install it by running `install.packages("dplyr")`. This is a one-time installation. Once per R session, you need to load the package using `library()`.

```
library(dplyr)
```

Now with `dplyr` loaded we can slice the protest data to just pick our certain rows, like the first row.

```
slice(dataProtest, 1)
```

	Date	Location	Attendees			
1	2017-01-15	Bowie State University, Bowie, MD	1500			
	Event..legacy..see.tags.				Tags Curated	
1		Healthcare Healthcare; For Affordable Care Act			Yes	
						Source
1		http://www.capitalgazette.com/news/ph-ac-cn-aca-rally-0116-20170115-story.html				
	Total.Articles					
1	1					

There is a more modern “grammar” in R using the pipe operator. This is a way to chain together functions in a more readable way. The pipe operator is `|>`. It takes the output of the function on the left and passes it as the first argument to the function on the right. Here is the same code as above using the pipe operator.

```
dataProtest |> slice(1)
```

	Date	Location	Attendees			
1	2017-01-15	Bowie State University, Bowie, MD	1500			
	Event..legacy..see.tags.				Tags Curated	
1		Healthcare Healthcare; For Affordable Care Act			Yes	
						Source
1		http://www.capitalgazette.com/news/ph-ac-cn-aca-rally-0116-20170115-story.html				
	Total.Articles					
1	1					

This code takes `dataProtest` and passes it in to the first argument of the `slice()` function. The `slice()` function then returns the first row of the dataset. The code is more readable this way.

You will also see many users using `%>%` in their code. The `%>%` pipe operator has been around longer, but the newer `|>` pipe operator, created in 2021 for R 4.1.0, is [faster](#). You can use either one.

If you want the first 3 rows you can also use `slice()`

```
dataProtest |> slice(1:3)
```

	Date	Location	Attendees
1	2017-01-15	Bowie State University, Bowie, MD	1500
2	2017-01-16	Johnson City, TN	300
3	2017-01-16	Indianapolis, IN	20

	Event...legacy...see.tags.
1	Healthcare
2	Civil Rights
3	Environment

	Tags	Curated
1	Healthcare; For Affordable Care Act	Yes
2	Civil Rights; For racial justice; Martin Luther King, Jr.	Yes
3	Environment; For wilderness preservation	Yes

1	http://www.capitalgazette.com/nac-cn-aca-rally-0116-20170115-story.html
2	http://www.johnsoncitypress.com/Local/2017/01/16/Hundreds-march-to-honor-Martin-Luther-King-Jr-in-Johnson-City.html?ci=featured&lp=46&ti=
3	http://wishtv.com/2017/01/16/nature-groups-protest-deforestation-of-crown-hill-cemetery-land/

	Total.Articles
1	1
2	4
3	1

or you can use `head()` that we used earlier.

```
dataProtest |> head(3)
```

	Date	Location	Attendees
1	2017-01-15	Bowie State University, Bowie, MD	1500

2	2017-01-16	Johnson City, TN	300
3	2017-01-16	Indianapolis, IN	20

Event..legacy..see.tags.

1	Healthcare
2	Civil Rights
3	Environment

	Tags Curated
1	Healthcare; For Affordable Care Act Yes
2	Civil Rights; For racial justice; Martin Luther King, Jr. Yes
3	Environment; For wilderness preservation Yes

1 <http://www.capitalgazette.com/nac-cn-aca-rally-0116-20170115-story.html>

2 <http://www.johnsoncitypress.com/Local/2017/01/16/Hundreds-march-to-honor-Martin-Luther-King-Jr-in-Johnson-City.html?ci=featured&lp=46&ti=>

3 <http://wishtv.com/2017/01/16/nature-groups-protest-deforestation-of-crown-hill-cemetery-land/>

Total.Articles	
1	1
2	4
3	1

I have the general habit of running `head()` and `tail()` on any datasets I am working with just to make sure it looks like what I expect. I encourage you to do the same. Many errors can be avoided by just looking at the data.

We may also be interested in only a few columns of the dataset. We can use the `select()` function to pick out the columns we want. For example, if we only want the date and location of the protest, we can use the following code.

```
dataProtest |>
  select(Date, Location) |>
  head(3)
```

	Date	Location
1	2017-01-15	Bowie State University, Bowie, MD
2	2017-01-16	Johnson City, TN
3	2017-01-16	Indianapolis, IN

This code takes `dataProtest` and passes it to the `select()` function. The `select()` function then returns only the `Date` and `Location` columns of the dataset. `head(3)` then returns the first 3 rows of the dataset. Here you can see how the pipe operator can be used to chain together functions in a readable way. Technically, this code is identical to

```
head(select(dataProtest, Date, Location), 3)
```

	Date	Location
1	2017-01-15	Bowie State University, Bowie, MD
2	2017-01-16	Johnson City, TN
3	2017-01-16	Indianapolis, IN

The computer does not care which approach you take. However, the potential problem with this code is that there is so much distance between `head` and the 3 at the end. This distance makes it harder to read, understand, and find errors. It will become even more important when we chain many more functions together.

You can also get a column by name using the `$` operator. For example, to get the `Date` column you can use `dataProtest$Date`. To get the first 10 dates you can use `dataProtest$Date[1:10]`. To get the first 10 locations you can use `dataProtest$Location[1:10]`.

```
dataProtest$Date[1:10]
```

```
[1] "2017-01-15" "2017-01-16" "2017-01-16" "2017-01-16" "2017-01-18"
[6] "2017-01-19" "2017-01-19" "2017-01-20" "2017-01-20" "2017-01-20"
```

```
dataProtest$Location[1:10]
```

```
[1] "Bowie State University, Bowie, MD"
[2] "Johnson City, TN"
[3] "Indianapolis, IN"
[4] "Cincinnati, OH"
[5] "Hartford, CT"
[6] "Washington, DC"
[7] "Washington, DC"
[8] "University of Washington, Seattle, WA"
[9] "Westlake Park, Seattle, WA"
[10] "Columbus, OH"
```

7 Indexing basics

The square brackets `[]` are used to pick out parts of an object by position. This is called indexing. R starts counting at 1 (I note this because some programming languages start at 0).

```
dataProtest$Date[1]
```

```
[1] "2017-01-15"
```

```
dataProtest$Date[1:5]
```

```
[1] "2017-01-15" "2017-01-16" "2017-01-16" "2017-01-16" "2017-01-18"
```

For a vector, like `dataProtest$Date`, we use one position inside the square brackets. For a data frame, which has rows and columns, we use two positions: `[row, column]`. The row position comes first, then a comma, then the column position.

```
dataProtest[1, 1]
```

```
[1] "2017-01-15"
```

```
dataProtest[1, ]
```

	Date	Location	Attendees	
1	2017-01-15	Bowie State University, Bowie, MD	1500	
	Event..legacy..see.tags.			Tags Curated
1		Healthcare Healthcare; For Affordable Care Act		Yes
				Source
1		http://www.capitalgazette.com/news/ph-ac-cn-aca-rally-0116-20170115-story.html		
	Total.Articles			
1			1	

```
dataProtest[, 1]
```

```
[1] "2017-01-15" "2017-01-16" "2017-01-16" "2017-01-16" "2017-01-18"
[6] "2017-01-19" "2017-01-19" "2017-01-20" "2017-01-20" "2017-01-20"
[11] "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-20"
[16] "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-20"
[21] "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-20"
[26] "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-20" "2017-01-21"
[31] "2017-01-21" "2017-01-21" "2017-01-21" "2017-01-21" "2017-01-21"
[36] "2017-01-21" "2017-01-21" "2017-01-21" "2017-01-21" "2017-01-21"
[41] "2017-01-21" "2017-01-21" "2017-01-21" "2017-01-21" "2017-01-21"
```


[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[15526]	"2018-12-04"	"2018-12-04"	"2018-12-04"	"2018-12-04"	"2018-12-04"
[15531]	"2018-12-05"	"2018-12-05"	"2018-12-05"	"2018-12-05"	"2018-12-05"
[15536]	"2018-12-05"	"2018-12-06"	"2018-12-06"	"2018-12-06"	"2018-12-06"
[15541]	"2018-12-06"	"2018-12-06"	"2018-12-06"	"2018-12-06"	"2018-12-06"
[15546]	"2018-12-06"	"2018-12-07"	"2018-12-07"	"2018-12-07"	"2018-12-07"
[15551]	"2018-12-07"	"2018-12-07"	"2018-12-07"	"2018-12-08"	"2018-12-08"
[15556]	"2018-12-08"	"2018-12-08"	"2018-12-08"	"2018-12-08"	"2018-12-09"
[15561]	"2018-12-09"	"2018-12-09"	"2018-12-10"	"2018-12-10"	"2018-12-10"
[15566]	"2018-12-10"	"2018-12-10"	"2018-12-10"	"2018-12-10"	"2018-12-10"
[15571]	"2018-12-10"	"2018-12-10"	"2018-12-10"	"2018-12-10"	"2018-12-11"
[15576]	"2018-12-11"	"2018-12-11"	"2018-12-12"	"2018-12-12"	"2018-12-12"
[15581]	"2018-12-12"	"2018-12-12"	"2018-12-12"	"2018-12-12"	"2018-12-12"
[15586]	"2018-12-12"	"2018-12-12"	"2018-12-12"	"2018-12-12"	"2018-12-12"
[15591]	"2018-12-12"	"2018-12-13"	"2018-12-13"	"2018-12-13"	"2018-12-14"
[15596]	"2018-12-14"	"2018-12-14"	"2018-12-14"	"2018-12-14"	"2018-12-14"
[15601]	"2018-12-15"	"2018-12-15"	"2018-12-15"	"2018-12-15"	"2018-12-15"
[15606]	"2018-12-15"	"2018-12-15"	"2018-12-15"	"2018-12-15"	"2018-12-15"
[15611]	"2018-12-15"	"2018-12-15"	"2018-12-15"	"2018-12-15"	"2018-12-16"
[15616]	"2018-12-16"	"2018-12-16"	"2018-12-16"	"2018-12-16"	"2018-12-16"
[15621]	"2018-12-17"	"2018-12-17"	"2018-12-17"	"2018-12-17"	"2018-12-17"
[15626]	"2018-12-17"	"2018-12-17"	"2018-12-17"	"2018-12-18"	"2018-12-18"
[15631]	"2018-12-18"	"2018-12-18"	"2018-12-18"	"2018-12-18"	"2018-12-19"
[15636]	"2018-12-19"	"2018-12-19"	"2018-12-19"	"2018-12-19"	"2018-12-19"
[15641]	"2018-12-19"	"2018-12-19"	"2018-12-19"	"2018-12-19"	"2018-12-20"
[15646]	"2018-12-20"	"2018-12-20"	"2018-12-20"	"2018-12-20"	"2018-12-21"
[15651]	"2018-12-21"	"2018-12-21"	"2018-12-21"	"2018-12-21"	"2018-12-21"
[15656]	"2018-12-22"	"2018-12-22"	"2018-12-22"	"2018-12-22"	"2018-12-22"
[15661]	"2018-12-23"	"2018-12-23"	"2018-12-23"	"2018-12-24"	"2018-12-24"
[15666]	"2018-12-25"	"2018-12-25"	"2018-12-27"	"2018-12-28"	"2018-12-28"
[15671]	"2018-12-28"	"2018-12-28"	"2018-12-28"	"2018-12-29"	"2018-12-29"
[15676]	"2018-12-29"	"2018-12-29"	"2018-12-29"	"2018-12-29"	"2018-12-30"
[15681]	"2018-12-30"	"2018-12-30"	"2018-12-31"	"2019-01-01"	"2019-01-02"
[15686]	"2019-01-02"	"2019-01-02"	"2019-01-02"	"2019-01-03"	"2019-01-03"
[15691]	"2019-01-03"	"2019-01-03"	"2019-01-03"	"2019-01-03"	"2019-01-03"
[15696]	"2019-01-03"	"2019-01-03"	"2019-01-03"	"2019-01-03"	"2019-01-03"
[15701]	"2019-01-03"	"2019-01-04"	"2019-01-04"	"2019-01-04"	"2019-01-05"
[15706]	"2019-01-05"	"2019-01-05"	"2019-01-05"	"2019-01-05"	"2019-01-05"
[15711]	"2019-01-05"	"2019-01-05"	"2019-01-05"	"2019-01-06"	"2019-01-06"
[15716]	"2019-01-06"	"2019-01-06"	"2019-01-06"	"2019-01-07"	"2019-01-07"
[15721]	"2019-01-07"	"2019-01-07"	"2019-01-07"	"2019-01-07"	"2019-01-07"
[15726]	"2019-01-07"	"2019-01-07"	"2019-01-08"	"2019-01-08"	"2019-01-08"
[15731]	"2019-01-08"	"2019-01-08"	"2019-01-08"	"2019-01-08"	"2019-01-08"
[15736]	"2019-01-08"	"2019-01-08"	"2019-01-09"	"2019-01-09"	"2019-01-09"

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[23481] "2020-03-06" "2020-03-06" "2020-03-06" "2020-03-07" "2020-03-07"
[23486] "2020-03-07" "2020-03-07" "2020-03-07" "2020-03-07" "2020-03-07"
[23491] "2020-03-07" "2020-03-07" "2020-03-07" "2020-03-07" "2020-03-07"
[23496] "2020-03-07" "2020-03-07" "2020-03-07" "2020-03-07" "2020-03-07"
[23501] "2020-03-08" "2020-03-08" "2020-03-08" "2020-03-09" "2020-03-09"
[23506] "2020-03-09" "2020-03-09" "2020-03-09" "2020-03-09" "2020-03-09"
[23511] "2020-03-09" "2020-03-09" "2020-03-09" "2020-03-09" "2020-03-09"
[23516] "2020-03-09" "2020-03-09" "2020-03-10" "2020-03-10" "2020-03-10"
[23521] "2020-03-10" "2020-03-10" "2020-03-10" "2020-03-10" "2020-03-10"
[23526] "2020-03-10" "2020-03-10" "2020-03-10" "2020-03-11" "2020-03-11"
[23531] "2020-03-11" "2020-03-11" "2020-03-11" "2020-03-11" "2020-03-11"
[23536] "2020-03-11" "2020-03-11" "2020-03-11" "2020-03-11" "2020-03-11"
[23541] "2020-03-11" "2020-03-11" "2020-03-12" "2020-03-12" "2020-03-13"
[23546] "2020-03-13" "2020-03-14" "2020-03-14" "2020-03-15" "2020-03-16"
[23551] "2020-03-18" "2020-03-19" "2020-03-20" "2020-03-21" "2020-03-22"
[23556] "2020-03-22" "2020-03-22" "2020-03-23" "2020-03-24" "2020-03-26"
[23561] "2020-03-26" "2020-03-26" "2020-03-27" "2020-03-27" "2020-03-27"
[23566] "2020-03-27" "2020-03-28" "2020-03-28" "2020-03-28" "2020-03-29"
[23571] "2020-03-29" "2020-03-29" "2020-03-30" "2020-03-30" "2020-03-30"
[23576] "2020-03-30" "2020-03-30" "2020-03-30" "2020-03-31" "2020-03-31"
[23581] "2020-03-31" "2020-04-01" "2020-04-01" "2020-04-01" "2020-04-01"
[23586] "2020-04-01" "2020-04-02" "2020-04-02" "2020-04-02" "2020-04-02"
[23591] "2020-04-02" "2020-04-02" "2020-04-02" "2020-04-02" "2020-04-02"
[23596] "2020-04-03" "2020-04-03" "2020-04-03" "2020-04-03" "2020-04-03"
[23601] "2020-04-03" "2020-04-03" "2020-04-03" "2020-04-04" "2020-04-04"
[23606] "2020-04-04" "2020-04-05" "2020-04-06" "2020-04-06" "2020-04-06"
[23611] "2020-04-06" "2020-04-06" "2020-04-06" "2020-04-06" "2020-04-07"
[23616] "2020-04-07" "2020-04-07" "2020-04-07" "2020-04-08" "2020-04-08"
[23621] "2020-04-08" "2020-04-08" "2020-04-08" "2020-04-08" "2020-04-08"
[23626] "2020-04-08" "2020-04-09" "2020-04-09" "2020-04-09" "2020-04-09"
[23631] "2020-04-09" "2020-04-09" "2020-04-09" "2020-04-10" "2020-04-10"
[23636] "2020-04-10" "2020-04-10" "2020-04-10" "2020-04-10" "2020-04-10"
[23641] "2020-04-10" "2020-04-10" "2020-04-11" "2020-04-11" "2020-04-11"
[23646] "2020-04-11" "2020-04-11" "2020-04-11" "2020-04-11" "2020-04-12"
[23651] "2020-04-12" "2020-04-12" "2020-04-12" "2020-04-12" "2020-04-13"
[23656] "2020-04-13" "2020-04-13" "2020-04-13" "2020-04-14" "2020-04-14"
[23661] "2020-04-14" "2020-04-14" "2020-04-14" "2020-04-14" "2020-04-15"
[23666] "2020-04-15" "2020-04-15" "2020-04-15" "2020-04-15" "2020-04-15"
[23671] "2020-04-15" "2020-04-15" "2020-04-15" "2020-04-15" "2020-04-15"
[23676] "2020-04-15" "2020-04-15" "2020-04-15" "2020-04-15" "2020-04-15"
[23681] "2020-04-15" "2020-04-15" "2020-04-16" "2020-04-16" "2020-04-16"
[23686] "2020-04-16" "2020-04-16" "2020-04-16" "2020-04-16" "2020-04-16"
[23691] "2020-04-16" "2020-04-16" "2020-04-16" "2020-04-16" "2020-04-16"

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

The code `dataProtest[1, 1]` returns the value in the first row and first column. The code `dataProtest[1, ]` returns the entire first row. The code `dataProtest[, 1]` returns the entire first column. Leaving one side of the comma blank means “give me all of them.”

You can also use column names inside the square brackets.

```
dataProtest[1:5, c("Date", "Location")]
```

	Date	Location
1	2017-01-15	Bowie State University, Bowie, MD
2	2017-01-16	Johnson City, TN
3	2017-01-16	Indianapolis, IN
4	2017-01-16	Cincinnati, OH
5	2017-01-18	Hartford, CT

In this script we will mostly use `slice()`, `select()`, `filter()`, and the `$` operator because they are often easier to read. But you will see bracket indexing in a lot of R code, so it is worth knowing the basic pattern.

8 Creating objects

So far every time we run some R code the results are dumped to the console. This is R’s default behavior. If you do not indicate otherwise, it will dump the results to the console and promptly forget those results. When we want to store the results, we can use the assignment operator `<-`. For example, to save the first 10 dates to a variable `a` you can use

```
a <- dataProtest$Date[1:10]
```

To save the first 10 locations to a variable `b` you can use

```
b <- dataProtest$Location[1:10]
```

Now if we run `ls()` we will see that we have two new variables `a` and `b` in our environment. We can use these variables later in our code.

```
ls()
```

```
[1] "a"          "b"          "dataProtest"
```

If you want to see the contents of a variable you can just type the variable name and run the code. For example, to see the contents of `a` you can run

```
a
```

```
[1] "2017-01-15" "2017-01-16" "2017-01-16" "2017-01-16" "2017-01-18"  
[6] "2017-01-19" "2017-01-19" "2017-01-20" "2017-01-20" "2017-01-20"
```

If a line of R code does not have a `<-`, then the results will not be stored. I would like to simplify our protest dataset by removing some columns that we will not use. I will use the `select()` function to pick out the columns to keep *and* use the `<-` operator to replace the original `dataProtest` with a new version of `dataProtest` that only has the columns I want.

```
dataProtest <- dataProtest |>  
  select(Date, Location, Attendees, Tags)
```

Now if you run `head(dataProtest)` you will see that the dataset only has the `Date`, `Location`, `Attendees`, and `Tags` columns. The other columns have been removed. `select()` also allows you to indicate which features to drop by prefixing their names with a minus sign. Instead of listing the features we wanted to keep, we could have listed the features we wanted to drop, using `select(-Event..legacy..see.tags., -Source, -Curated, -Total.Articles)`.

When naming objects in R, choose names that are short, clear, and meaningful. Object names can contain letters, numbers, underscores (`_`), and periods (`.`), but they should not start with a number and should not contain spaces. As previously noted, R is case sensitive, so `protestData`, `ProtestData`, and `protestdata` are three different objects. I like starting with a lowercase label of what the object is, like the `data` in `dataProtest`. I tend to use “camel case,” which capitalizes separate words in an object name, like `listIceCreamFlavors`. Clear names make your code easier to read, easier to debug, and easier to return to later when you no longer remember exactly what you were thinking.

8.1 Exercises

1. What is the date of the protest in line 10000 of the dataset?
2. Which protest type is in line 4289 of the dataset?

9 Filtering rows

We can ask every location if they equal “Philadelphia, PA”.

```
# let's just ask the first 10, otherwise will print out the first 1,000
dataProtest$Location[1:10]=="Philadelphia, PA"
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

Note the use of the double equal sign `==`. This is the “logical” equal. It is not making `Location` equal to Philadelphia, PA. It is asking if `Location` is equal to Philadelphia, PA. The result is a vector of `TRUE` and `FALSE` values. If the location is Philadelphia, PA, then the result is `TRUE`. If the location is not Philadelphia, PA, then the result is `FALSE`.

How many protests occurred in Philadelphia, PA?

```
dataProtest |>
  filter(Location=="Philadelphia, PA") |>
  nrow()
```

```
[1] 193
```

The `filter()` function is used to select rows that meet a certain condition. In this case, we are selecting rows where the `Location` is equal to “Philadelphia, PA”. The expression `Location=="Philadelphia, PA"` will evaluate to `TRUE` for any row where `Location` is identical to “Philadelphia, PA” and `FALSE` otherwise. `filter()` will keep only those rows where the logical expression evaluates to `TRUE` eliminating all others (NAs also get eliminated). The `nrow()` function, which we met earlier, is used to count the number of rows in the dataset. The result is the number of protests that occurred in Philadelphia, PA.

However, this count does not include those with locations like “University of Pennsylvania, Philadelphia, PA”. For example, these ones:

```
dataProtest |>
  filter(Location=="University of Pennsylvania, Philadelphia, PA")
```

	Date	Location	Attendees
1	2018-02-22	University of Pennsylvania, Philadelphia, PA	130
2	2019-04-23	University of Pennsylvania, Philadelphia, PA	10
3	2019-04-23	University of Pennsylvania, Philadelphia, PA	50
4	2019-10-23	University of Pennsylvania, Philadelphia, PA	NA

	Tags
1	Guns; For greater gun control
2	Other; For animal welfare
3	Other; Against closure/relocation
4	Immigration; For compassionate immigration; Against invited speaker

The `Location` feature has the phrase “Philadelphia, PA”, but the `Location` is not *exactly* identical to “Philadelphia, PA”. It is time to introduce you to `grepl()`, which is a very powerful function for searching for patterns in text. For now, we will use it simply to search for any `Location` containing the phrase “Philadelphia, PA”. `grepl()` returns `TRUE` if the phrase is found and `FALSE` if it is not found. For example, to find all protests that occurred in Philadelphia, PA, we can use the following code.

```
dataProtest |>
  filter(grepl("Philadelphia, PA", Location)) |>
  head(n=5)
```

	Date	Location	Attendees
1	2017-01-21	Philadelphia, PA	50000
2	2017-01-26	Philadelphia, PA	2360
3	2017-01-29	Philadelphia International Airport, Philadelphia, PA	1910
4	2017-02-02	Philadelphia, PA	800
5	2017-02-04	Philadelphia City Hall, Philadelphia, PA	2000

	Tags
1	Civil Rights; For women's rights; Women's March
2	Executive; Against 45th president
3	Immigration; Against travel ban
4	Immigration; Against travel ban
5	Immigration; Against travel ban

Now we have found many more protests in Philadelphia since some of them were at the airport or at City Hall. Let’s redo that count.

```
dataProtest |>
  filter(grepl("Philadelphia, PA", Location)) |>
  nrow()
```

```
[1] 327
```

We will study `grepl()` and its variants a lot more later, but for now think of it as “Find” in your word processor. If you are looking for a word in a document, you can use “Find” to locate all instances of that word. `grepl()` is the same idea. It is looking for a phrase in a text field.

We can include multiple conditions in the `filter()` function. For example, to find all protests in Philadelphia, PA, before 2018 with more than 1,000 attendees, we can use the following code. Note that `&` is the logical AND operator. It returns `TRUE` if both conditions are `TRUE` and `FALSE` otherwise. The `|` operator is the logical OR operator. It returns `TRUE` if either condition is `TRUE` and `FALSE` otherwise.

```
dataProtest |>
  filter(grepl("Philadelphia, PA", Location) &
         (Date <= "2017-12-31") &
         (Attendees >= 1000))
```

	Date	Location	Attendees
1	2017-01-21	Philadelphia, PA	50000
2	2017-01-26	Philadelphia, PA	2360
3	2017-01-29	Philadelphia International Airport, Philadelphia, PA	1910
4	2017-02-04	Philadelphia City Hall, Philadelphia, PA	2000
5	2017-03-02	Independence Mall, Philadelphia, PA	1000
6	2017-04-15	Philadelphia, PA	2000
7	2017-04-22	Philadelphia, PA	10000
8	2017-04-29	Philadelphia, PA	2000
9	2017-05-01	Philadelphia, PA	2000
10	2017-05-01	Philadelphia, PA	1000
11	2017-08-16	Philadelphia, PA	2000

	Tags
1	Civil Rights; For women's rights; Women's March
2	Executive; Against 45th president
3	Immigration; Against travel ban
4	Immigration; Against travel ban
5	Civil Rights; For religious tolerance
6	Executive; Against 45th president; Tax returns
7	Other; For science; March for Science
8	Environment; Against climate change; People's Climate March
9	Immigration; For compassionate immigration; For worker rights; May Day
10	Collective Bargaining; For better compensation; May Day
11	Civil Rights; For racial justice; Against white supremacy; Charlottesville

9.1 Exercise

- How many protests occurred in your home state? If not from the US just pick a state like New York “NY” or California “CA” or Pennsylvania “PA”
- Where did the protest in the last row of the full dataset occur?

10 Summarizing data

What is the average size of a protest? The `summarize()` function is used to calculate summary statistics. For example, to calculate the average number of attendees at a protest, we can use

the following code.

```
dataProtest |>
  summarize(mean(Attendees))
```

```
mean(Attendees)
1              NA
```

Hmmm... it looks like there are some missing values in the `Attendees` column. Rather than just dropping them and computing the average of the rest, R forces us to be intentional about handling NAs. If indeed we want to drop the NAs, then we can use the `na.rm=TRUE` argument to remove the missing values before calculating the average.

```
dataProtest |>
  summarize(mean(Attendees, na.rm=TRUE))
```

```
mean(Attendees, na.rm = TRUE)
1              643.8831
```

Perhaps we are interested in several data summaries at the same time. No problem. Just include them all in `summarize()`.

```
dataProtest |>
  summarize(average = mean(Attendees, na.rm=TRUE),
           median = median(Attendees, na.rm=TRUE),
           minimum = min(Attendees, na.rm=TRUE),
           maximum = max(Attendees, na.rm=TRUE),
           NAcount = sum(is.na(Attendees)))
```

```
average median minimum maximum NAcount
1 643.8831    100        0  725000    15061
```

That was a lot of typing to get a complete set of summary statistics. The `summary()` function is always available for that.

```
summary(dataProtest$Attendees)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NAs
0.00	26.75	100.00	643.88	200.00	725000.00	15061

You can also use it to get a quick summary of the entire dataset.

```
summary(dataProtest)
```

Date	Location	Attendees	Tags
Length :38097	Length :38097	Min. : 0.00	Length :38097
N.unique : 1464	N.unique :10037	1st Qu.: 26.75	N.unique : 2208
N.blank : 0	N.blank : 0	Median : 100.00	N.blank : 0
Min.nchar: 10	Min.nchar: 5	Mean : 643.88	Min.nchar: 14
Max.nchar: 10	Max.nchar: 75	3rd Qu.: 200.00	Max.nchar: 142
		Max. :725000.00	
		NA's :15061	

11 Mutate to edit and create new columns

The data does not contain a column for the state in which the protest occurred. We can create this column by extracting the state from the `Location` column. The last two characters of the `Location` column contain the state abbreviation. We can use the `str_sub()` function from the `stringr` package to extract the last two characters of the `Location` column. The `str_sub()` function is used to extract a substring from a string. For example, to extract the last two characters of the string “Philadelphia, PA”, we can use the following code. Let’s load the `stringr` and test out `str_sub()` on an example.

```
library(stringr)
str_sub("Philadelphia, PA", -2)
```

```
[1] "PA"
```

The first argument is the string from which to extract the substring. The second argument is the starting position of the substring. A nice feature of `str_sub()` is that you can use negative numbers which it interprets as characters from the end. So the -2 tells `str_sub()` to start at the second to last character. The third argument is the ending position of the substring. Here the -1 means the very last character of the string. If we do not include a third argument, then `str_sub()` will extract the substring starting at the second argument and continuing to the end of the string.

```
str_sub("Philadelphia, PA", -2)
```

```
[1] "PA"
```


There are other R functions that can extract substrings including `substring()`, `substr()`, and `gsub()`. I am introducing you to `str_sub()` because it is the only one that lets you put negative numbers in the second and third arguments to easily grab substrings from the end. This is a very useful feature.

With `str_sub()` now in our toolbox, we can make a new column called `state` that contains the state in which the protest occurred.

```
dataProtest <- dataProtest |>
  mutate(state=str_sub(Location, -2))
head(dataProtest)
```

	Date	Location	Attendees		Tags	state
1	2017-01-15	Bowie State University, Bowie, MD	1500			
2	2017-01-16	Johnson City, TN	300			
3	2017-01-16	Indianapolis, IN	20			
4	2017-01-16	Cincinnati, OH	NA			
5	2017-01-18	Hartford, CT	300			
6	2017-01-19	Washington, DC	NA			
1		Healthcare; For Affordable Care Act				MD
2		Civil Rights; For racial justice; Martin Luther King, Jr.				TN
3		Environment; For wilderness preservation				IN
4		Civil Rights; For racial justice; Martin Luther King, Jr.				OH
5		Healthcare; For Planned Parenthood				CT
6		Executive; Against 45th president				DC

Peeking at the first few rows of `dataProtest` we can see that there is a new column with the state abbreviation. Please always check that your code does what you intended to do. Run, check, run, check, one line at a time.

So you can see that `mutate()` is useful for making new data features computed based on other features. We also will use it to edit or clean up data. Let's check what these state abbreviations look like.

```
dataProtest |>
  count(state)
```

	state	n
1	AK	252
2	AL	281
3	AR	174

4	AZ	563
5	CA	4439
6	CO	813
7	CT	708
8	DC	536
9	DE	115
10	FL	1822
11	Fl	1
12	GA	623
13	GU	22
14	HI	182
15	Hi	1
16	IA	470
17	ID	344
18	IL	1273
19	IN	700
20	KS	293
21	KY	821
22	LA	330
23	MA	1265
24	MD	453
25	ME	437
26	MI	1410
27	MN	747
28	MO	800
29	MS	187
30	MT	294
31	Mi	1
32	NC	1150
33	ND	98
34	NE	257
35	NH	266
36	NJ	893
37	NM	402
38	NV	300
39	NY	2688
40	OH	1107
41	OK	324
42	OR	1368
43	PA	1656
44	PR	19
45	RI	194
46	SC	439

```

47 SD 101
48 TN 576
49 TX 1649
50 UT 421
51 VA 906
52 VT 337
53 WA 1375
54 WI 812
55 WV 266
56 WY 131
57 ce 1
58 co 1
59 iD 1
60 te 1
61 wA 1

```

Here I have used the `count()` function to count the number of protests in each state. It groups the data by the `state` column and then counts the number of rows in each group. The result is a new data frame with one column containing the state abbreviation (`state`) and another column containing the number of protests in that state (`count()` will always call this one `n`).

Do you see some problems with our state abbreviations? I see an “Fl”, an “Hi”, and an “Mi” and a few others that do not seem to be correctly capitalized. I also see some abbreviations that are “CE” and “TE”, not states that I know of. Let’s take a closer look at these strange ones. Note that I am introducing the `%in%` operator. This is a logical operator that asks each value of `state` whether its value is in the collection to the right of `%in%`. It is a more compact way to write `state=="Fl" | state=="Hi" | state=="Mi" | state=="ce" | state=="co" | state=="iD" | state=="te" | state=="wA"`. Well, there. I have gone ahead and typed that all out. I hope to never have to type a logical expression with so many ORs again.

```

dataProtest |>
  filter(state %in% c("Fl","Hi","Mi","ce","co","iD","te","wA")) |>
  select(state, Location)

```

	state	Location
1	co	Ciudad Juarez, Mexico
2	ce	Space
3	Fl	Panama City, Fl
4	Mi Wyoming	Godfrey-Lee High School, Wyoming, Mi
5	Hi	Honolulu, Hi
6	wA	Montesano, wA
7	iD	City Hall, Pocatello, iD
8	te	La Porte County Courthouse in La Porte

Perhaps even more straightforward, R has a built in list of state abbreviations, `state.abb`. We can just filter those values of `state` that are not in this list (I will add Washington DC, Puerto Rico, and Guam too).

```
dataProtest |>
  filter(!(state %in% c(state.abb, "DC", "PR", "GU"))) |>
  select(state, Location)
```

	state	Location
1	co	Ciudad Juarez, Mexico
2	ce	Space
3	Fl	Panama City, Fl
4	Mi Wyoming	Godfrey-Lee High School, Wyoming, Mi
5	Hi	Honolulu, Hi
6	wA	Montesano, wA
7	iD	City Hall, Pocatello, iD
8	te	La Porte County Courthouse in La Porte

Lots of different kinds of errors here. Five of them are just lower case. One is in Mexico (we need to drop this one). One is in Space (space is cool so let's keep that one for fun), and one is in La Porte, which I had to look up La Porte to find that it is in Indiana (IN). Let's clean this up using `mutate()`.

```
dataProtest <- dataProtest |>
  filter(state != "co") |> # drop Mexico
  mutate(state = toupper(state),
         state = replace_values(state,
                                "CE" ~ "Space",
                                "TE" ~ "IN"))

dataProtest |>
  count(state)
```

	state	n
1	AK	252
2	AL	281
3	AR	174
4	AZ	563
5	CA	4439
6	CO	813
7	CT	708
8	DC	536

9	DE	115
10	FL	1823
11	GA	623
12	GU	22
13	HI	183
14	IA	470
15	ID	345
16	IL	1273
17	IN	701
18	KS	293
19	KY	821
20	LA	330
21	MA	1265
22	MD	453
23	ME	437
24	MI	1411
25	MN	747
26	MO	800
27	MS	187
28	MT	294
29	NC	1150
30	ND	98
31	NE	257
32	NH	266
33	NJ	893
34	NM	402
35	NV	300
36	NY	2688
37	OH	1107
38	OK	324
39	OR	1368
40	PA	1656
41	PR	19
42	RI	194
43	SC	439
44	SD	101
45	Space	1
46	TN	576
47	TX	1649
48	UT	421
49	VA	906
50	VT	337
51	WA	1376

```
52    WI    812
53    WV    266
54    WY    131
```

Several things are happening here. First, the `toupper()` function is used to make sure that all state abbreviations are in upper case. Then we are using `replace_values()` to change the state abbreviations that need special fixes. Note its structure. The first argument is the variable that we are changing (`state`). Then we list all the replacements that we want to make. We are changing “CE” to “Space” and “TE” to “IN”. Finally we rerun the `count()` function to see if our changes worked. All looks good now.

The last feature that we have yet to explore is the `Tags` column. This column contains a list of reasons for the protest. The format of the tags is to have the reasons separated by a semicolon and a space. For example, a protest might have the tags “Civil Rights; Against pandemic intervention; Police brutality”. We can use the `strsplit()` function to split the tags into separate reasons. For example, to split the tags in the first three rows of the dataset, we can use the following code.

```
# what does the tag look like originally?
dataProtest$Tags[1:3]
```

```
[1] "Healthcare; For Affordable Care Act"
[2] "Civil Rights; For racial justice; Martin Luther King, Jr."
[3] "Environment; For wilderness preservation"
```

```
# now split it
strsplit(dataProtest$Tags[1:3], "; ")
```

```
[[1]]
[1] "Healthcare"          "For Affordable Care Act"

[[2]]
[1] "Civil Rights"        "For racial justice"
[3] "Martin Luther King, Jr."

[[3]]
[1] "Environment"         "For wilderness preservation"
```

`strsplit()` returns a `list` structure. This is a structure in R that has no columns and rows. Since each protest has a different number of tags, once we split them up, they do not fit neatly into fixed columns. We can use `unlist()` to remove the list structure and create a long vector

of all of the tags. And I will use `table()`, `sort()`, and `tail()` to find the most common reasons for a protest.

```
reasons <- strsplit(dataProtest$Tags, "; ")
reasons <- unlist(reasons)
table(reasons) |> sort() |> tail()
```

```
reasons
```

	Immigration	Other
	3543	4556
	Police For greater accountability	
	8254	8376
	For racial justice	Civil Rights
	10575	14807

Clearly, Civil Rights has topped the list. We can use this information to create a new column that is 1 if the protest has the tag “Civil Rights” and 0 otherwise.

```
dataProtest <- dataProtest |>
  mutate(civilrights = as.numeric(grepl("Civil Rights", Tags)))
```

Just like before when we used `grepl()` to find any text matches for “Philadelphia, PA”, this time we are using it to search `Tags` for any matches to “Civil Rights”. Again, it returns `TRUE` if the pattern is found and `FALSE` otherwise. `as.numeric()` converts `TRUE` to 1 and `FALSE` to 0.

This script is getting long. I have done every step piece by piece with a lot of explanation in between. In practice, you would not do this. You would combine everything into one pipeline that takes in the original dataset and does all the filtering and mutating and selecting to get you the dataset that you want. Here is everything we have done so far compactly written.

```
load("protests.RData")
dataProtest <- dataProtest |>
  select(Date, Location, Attendees, Tags) |>
  filter(Location != "Ciudad Juarez, Mexico") |>
  mutate(state = str_sub(Location, -2),
         state = toupper(state),
         state = replace_values(state,
                                "CE" ~ "Space",
                                "TE" ~ "IN"),
         civilrights = as.numeric(grepl("Civil Rights", Tags)))
head(dataProtest)
```

	Date	Location	Attendees			
1	2017-01-15	Bowie State University, Bowie, MD	1500			
2	2017-01-16	Johnson City, TN	300			
3	2017-01-16	Indianapolis, IN	20			
4	2017-01-16	Cincinnati, OH	NA			
5	2017-01-18	Hartford, CT	300			
6	2017-01-19	Washington, DC	NA			
				Tags	state	civilrights
1		Healthcare; For Affordable Care Act		MD		0
2	Civil Rights; For racial justice; Martin Luther King, Jr.			TN		1
3	Environment; For wilderness preservation			IN		0
4	Civil Rights; For racial justice; Martin Luther King, Jr.			OH		1
5	Healthcare; For Planned Parenthood			CT		0
6	Executive; Against 45th president			DC		0

11.1 Exercises

5. Which state had the most protests?
6. Which state had the least protests?
7. Which state had the most civil rights protests?
8. Create a new column that is 1 if the protest has the tag 'Against pandemic intervention'
9. Which state had the most protests against pandemic interventions?

12 Creating your own functions

Part of what makes R so powerful and useful is that you can create your own functions. In this way, the R user community can expand R's capabilities to do new tasks. For example, R does not have a built-in function to find the most common value in a collection. We can create our own function to do this. Have a look at this sequence of steps.

```
a <- table(unlist(reasons))
a |> head()
```

Against 45th president	Against 46th president
1543	3
Against abortion rights	Against accusations
444	2


```
Against administrative leave
                        6
```

```
Against advisor
                        12
```

```
max(a)
```

```
[1] 14807
```

```
a[a==max(a)]
```

```
Civil Rights
      14807
```

```
names(a[a==max(a)])
```

```
[1] "Civil Rights"
```

You have seen `table()` and `unlist()` in action earlier. Then I used `max()` to find the largest number of protests for a single reason. Then I used the expression `a[a==max(a)]`. Inside the square brackets, I ask each value of `a` (the table counts) if they equal the largest value. This returns a logical vector of `TRUE` and `FALSE` values. The square brackets will then pick out from `a` only those values where the logical expression `a==max(a)` evaluates to `TRUE`. I use this approach rather than `max()` or `head(1)` because it is possible that there are multiple tags that equal the maximum count. Finally, I used `names()` to get the name of the reason. I can pack all of this into a new function called `mostCommon()`.

```
mostCommon <- function(x)
{
  a <- table(x)
  return( names(a[a==max(a)]) )
}
```

This function is now a part of our R session and we can use it as we have other functions like `max()` or `mean()`. For example, to find the state with the most protests:

```
mostCommon(dataProtest$state)
```

```
[1] "CA"
```

Or the most common date for a protest.

```
mostCommon(dataProtest$Date)
```

```
[1] "2018-03-14"
```

What the most common date for civil rights protests in Texas?

```
dataProtest |>
  filter(state=="TX" & civilrights==1) |>
  summarize(mostCommon(Date))
```

```
mostCommon(Date)
1      2020-06-06
```

What happened in Texas on 2020-06-06?

```
dataProtest |>
  filter(Date=="2020-06-06" & state=="TX") |>
  count(Tags)
```

	Tags	n
1	Civil Rights; For racial justice; For greater accountability; Police	28
2	Civil Rights; For white supremacy; Counter protest	1
3	Guns; Against greater gun control	1

This is the height of the George Floyd protests. There were 28 protests recorded in Texas on that day tagged with “Civil Rights; For racial justice; For greater accountability; Police”.

Let’s make a special collection of states that includes PA and all of its bordering states. We can use this collection to filter the dataset to only include protests in these states.

```
PAplusBorderingstates <- c("PA","DE","MD","NJ","NY","OH","WV")
dataProtest |>
  filter(state %in% PAplusBorderingstates) |>
  summarize(mostCommon(Date))
```

```
mostCommon(Date)
1      2018-03-14
```

As I did earlier, I used the `%in%` operator to ask each state in `dataProtest` whether it is a member of the `PAplusBorderingstates` collection. This returns a logical vector of `TRUE` and `FALSE` values. The `filter()` function then keeps only those rows where the logical expression evaluates to `TRUE`.

Here we find that 2018-03-14 is the most common date for protests in Pennsylvania and its bordering states. This particular pi-Day was the day of the National School Walkout to protest gun violence.

```
dataProtest |>
  filter(Date=="2018-03-14" & state %in% PAplusBorderingstates) |>
  count(Tags)
```

	Tags	n
1	Civil Rights; For freedom of speech	1
2	Civil Rights; For racial justice; For greater accountability; Police	1
3	Environment; Against fossil fuels	1
4	Guns; Against greater gun control; Counter protest	2
5	Guns; For greater gun control	2
6	Guns; For greater gun control; National Walkout Day	262

13 Summarizing with groups of protests

We can use the `group_by()` function to group the data by a certain feature. All subsequent operations will be performed separately within each group. For example, let's total the number of protest attendees by state.

```
# will double count protesters at multiple protests
dataProtest |>
  group_by(state) |>
  summarize(sum(Attendees, na.rm=TRUE)) |>
  print(n=Inf)
```

```
# A tibble: 54 x 2
  state `sum(Attendees, na.rm = TRUE)`
  <chr> <int>
1 AK   35987
2 AL   34919
3 AR   21859
4 AZ  224194
```

5	CA	3190858
6	CO	428654
7	CT	106285
8	DC	1460536
9	DE	11280
10	FL	413328
11	GA	177400
12	GU	945
13	HI	65548
14	IA	101200
15	ID	45776
16	IL	907239
17	IN	95985
18	KS	45736
19	KY	111992
20	LA	45151
21	MA	507235
22	MD	70662
23	ME	80716
24	MI	214651
25	MN	253084
26	MO	130153
27	MS	21677
28	MT	66652
29	NC	230558
30	ND	13599
31	NE	72351
32	NH	45947
33	NJ	166706
34	NM	88496
35	NV	95383
36	NY	1730569
37	OH	182713
38	OK	74817
39	OR	393032
40	PA	391832
41	PR	15420
42	RI	35288
43	SC	71799
44	SD	16353
45	Space	0
46	TN	166575
47	TX	1136339

48	UT	93693
49	VA	127368
50	VT	68376
51	WA	490261
52	WI	211482
53	WV	31804
54	WY	11929

`summarize()` calculated the total number of attendees within each state. By default, R will print only the first 10 rows of the dataset. I used `print(n=Inf)` to force R to print all the rows.

We can also calculate the average number of attendees at a protest in each state.

```
options(pillar.sigfig=5) # less rounding
dataProtest |>
  group_by(state) |>
  summarize(Total=sum(Attendees, na.rm=TRUE),
            Average=mean(Attendees, na.rm=TRUE)) |>
  print(n=Inf)
```

```
# A tibble: 54 x 3
  state   Total Average
  <chr>   <int>   <dbl>
1 AK      35987   218.10
2 AL      34919   231.25
3 AR       21859   208.18
4 AZ     224194   640.55
5 CA    3190858  1191.1
6 CO     428654   865.97
7 CT     106285   238.84
8 DC    1460536  4651.4
9 DE       11280   163.48
10 FL     413328   382.36
11 GA     177400   476.88
12 GU        945     63
13 HI      65548   550.82
14 IA     101200   328.57
15 ID      45776   293.44
16 IL     907239  1154.2
17 IN      95985   195.49
18 KS      45736   245.89
```

19	KY	111992	288.64
20	LA	45151	226.89
21	MA	507235	604.57
22	MD	70662	245.35
23	ME	80716	271.77
24	MI	214651	257.99
25	MN	253084	562.41
26	MO	130153	309.15
27	MS	21677	216.77
28	MT	66652	320.44
29	NC	230558	347.75
30	ND	13599	209.22
31	NE	72351	411.09
32	NH	45947	268.70
33	NJ	166706	289.92
34	NM	88496	330.21
35	NV	95383	456.38
36	NY	1730569	1070.2
37	OH	182713	295.65
38	OK	74817	413.35
39	OR	393032	517.15
40	PA	391832	352.05
41	PR	15420	1401.8
42	RI	35288	273.55
43	SC	71799	276.15
44	SD	16353	247.77
45	Space	0	NaN
46	TN	166575	470.55
47	TX	1136339	1228.5
48	UT	93693	331.07
49	VA	127368	225.43
50	VT	68376	309.39
51	WA	490261	604.51
52	WI	211482	473.11
53	WV	31804	200.03
54	WY	11929	151

I used `options(pillar.sigfig=5)` to show more digits of precision in the output.

Interested in which “state” has the largest average protest size? Use `slice_max()`.

```
dataProtest |>
  group_by(state) |>
  summarize(Average=mean(Attendees, na.rm=TRUE)) |>
  slice_max(Average)
```

```
# A tibble: 1 x 2
  state Average
  <chr>   <dbl>
1 DC     4651.4
```

We can also simply arrange the rows in descending order of average protest size.

```
dataProtest |>
  group_by(state) |>
  summarize(Average=mean(Attendees, na.rm=TRUE)) |>
  arrange(desc(Average))
```

```
# A tibble: 54 x 2
  state Average
  <chr>   <dbl>
1 DC     4651.4
2 PR     1401.8
3 TX     1228.5
4 CA     1191.1
5 IL     1154.2
6 NY     1070.2
7 CO      865.97
8 AZ      640.55
9 MA      604.57
10 WA      604.51
# i 44 more rows
```

13.1 Exercises

10. Are civil rights protests larger on average than non-civil rights protests? (Hint: use `group_by/summarize`)

14 Pivot wider and longer

Note: You may encounter code with `melt()`, `cast()`, `reshape()`, `gather()`, and `spread()`. All of these are legacy versions of `pivot_wider()` and `pivot_longer()`.

`pivot_wider()` and `pivot_longer()` reorganize datasets between “wide” and “long” form. `pivot_longer()` takes many side-by-side columns (say, one column per year) and stacks them into two neat columns, one that says “which year” and one that shows the value, so your table becomes taller and easier to plot or compare. `pivot_wider()` does the opposite. It takes a tall list of items (like many rows per state and year) and spreads them back out so each thing gets its own column, making the table wider and easier to read at a glance.

`pivot_wider()` and `pivot_longer()` are in the `tidyr` package.

```
library(tidyr)
```

Let’s say that we are interested in determining which states have fewer civil rights protesters than non-civil rights protesters. With our existing `group_by()` and `summarize()` skills we can tabulate the number of protesters by state and by type of protest.

```
dataProtest |>
  group_by(state, civilrights) |>
  summarize(totAttendees = sum(Attendees, na.rm=TRUE))
```

```
`summarise()` has regrouped the output.
i Summaries were computed grouped by state and civilrights.
i Output is grouped by state.
i Use `summarise(.groups = "drop_last")` to silence this message.
i Use `summarise(.by = c(state, civilrights))` for per-operation grouping
  (`?dplyr::dplyr_by`) instead.
```

```
# A tibble: 107 x 3
# Groups:   state [54]
  state civilrights totAttendees
  <chr>      <dbl>         <int>
1 AK          0           15512
2 AK          1           20475
3 AL          0           11321
4 AL          1           23598
5 AR          0            8902
6 AR          1           12957
7 AZ          0          124949
```



```

8 AZ          1      99245
9 CA          0      701194
10 CA         1     2489664
# i 97 more rows

```

Now we can see the protester counts for civil rights and non-civil rights protest for each state. However, it is not clear how to compare the rows within each state.

This is where `pivot_wider()` comes in handy. We tell `pivot_wider()` to take the `civilrights` column and spread its values across columns, creating a column for non-civil rights protester counts and another column for civil rights protester counts. The values that will fill the new table come from `totAttendees`. Since `civilrights` takes values 0 and 1, I have asked `pivot_wider()` to paste “CR” in front of the values so that we get valid R column names.

```

dataProtest |>
  group_by(state, civilrights) |>
  summarize(totAttendees = sum(Attendees, na.rm=TRUE)) |>
  pivot_wider(names_from = civilrights,
              values_from = totAttendees,
              names_prefix = "CR")

```

```

`summarise()` has regrouped the output.
i Summaries were computed grouped by state and civilrights.
i Output is grouped by state.
i Use `summarise(.groups = "drop_last")` to silence this message.
i Use `summarise(.by = c(state, civilrights))` for per-operation grouping
  (`?dplyr::dplyr_by`) instead.

```

```

# A tibble: 54 x 3
# Groups:   state [54]
   state    CR0    CR1
   <chr> <int> <int>
1 AK     15512  20475
2 AL     11321  23598
3 AR      8902  12957
4 AZ    124949  99245
5 CA    701194 2489664
6 CO     88782  339872
7 CT     52659  53626
8 DC    477601  982935
9 DE       5507   5773
10 FL    222991 190337
# i 44 more rows

```

Almost done. We just need to filter those with fewer civil rights protesters.

```
dataProtest |>
  group_by(state, civilrights) |>
  summarize(totAttendees = sum(Attendees, na.rm=TRUE)) |>
  pivot_wider(names_from = civilrights,
              values_from = totAttendees,
              names_prefix = "CR") |>
  filter(CR1 < CR0)
```

`summarise()` has regrouped the output.
i Summaries were computed grouped by state and civilrights.
i Output is grouped by state.
i Use `summarise(.groups = "drop\_last")` to silence this message.
i Use `summarise(.by = c(state, civilrights))` for per-operation grouping
(`?dplyr::dplyr\_by`) instead.

```
# A tibble: 11 x 3
# Groups:   state [11]
  state    CR0    CR1
  <chr> <int> <int>
1 AZ    124949 99245
2 FL    222991 190337
3 GU      715   230
4 HI    36868 28680
5 ID    27156 18620
6 IN    49362 46623
7 MT    37871 28781
8 OK    50503 24314
9 PR    15320   100
10 VA    69230 58138
11 WV    24431  7373
```

Spreading the columns out wide made it easier to make our calculation within each state. There is usually a way to avoid `pivot_wider()` and still get the right answer. For example,

```
dataProtest |>
  group_by(state) |>
  summarize(CR0 = sum(if_else(civilrights == 0,
                             Attendees, 0, missing = 0), na.rm = TRUE),
           CR1 = sum(if_else(civilrights == 1,
```

```
Attendees, 0, missing = 0), na.rm = TRUE)) |>
filter(CR1 < CR0)
```

```
# A tibble: 11 x 3
  state    CR0    CR1
  <chr> <dbl> <dbl>
1 AZ    124949 99245
2 FL    222991 190337
3 GU      715   230
4 HI    36868 28680
5 ID    27156 18620
6 IN    49362 46623
7 MT    37871 28781
8 OK    50503 24314
9 PR    15320   100
10 VA    69230 58138
11 WV    24431  7373
```

This approach requires a little trick with the `if_else()`, including the value of `Attendees` in the sum if the associated protest is/is not a civil rights protest. Still, having `pivot_wider()` in your toolbox will come in handy.

Let's try an example that is a little more complicated. We will determine which states had the largest percent increase in protesters in 2020 compared to the average annual number of protesters 2017-2019. For each state we need a total for 2020 and an average annual number across 2017, 2018, and 2019.

First, let's derive a few helpful columns. Extract the year from `Date` and label that year as either 2017-2019 or as 2020. Note that I used `2017_2019` with an underscore as the label since R can confuse a hyphen for a minus sign.

```
dataProtest |>
  filter(Date < "2021-01-01") |>
  mutate(year = substring(Date, 1, 4) |> as.numeric(),
         timePeriod = ifelse(year %in% 2017:2019,
                              "2017_2019",
                              "2020")) |>
  select(Attendees, civilrights, Date, year, timePeriod) |>
  head()
```

```
Attendees civilrights    Date year timePeriod
1      1500           0 2017-01-15 2017 2017_2019
```

2	300	1	2017-01-16	2017	2017_2019
3	20	0	2017-01-16	2017	2017_2019
4	NA	1	2017-01-16	2017	2017_2019
5	300	0	2017-01-18	2017	2017_2019
6	NA	0	2017-01-19	2017	2017_2019

Second, we can get the total number of protesters by state and time period.

```
dataProtest |>
  filter(Date < "2021-01-01") |>
  mutate(year = substring(Date, 1, 4) |> as.numeric(),
         timePeriod = ifelse(year %in% 2017:2019,
                              "2017_2019",
                              "2020")) |>
  group_by(state, timePeriod) |>
  summarize(totProtesters = sum(Attendees, na.rm=TRUE))
```

`summarise()` has regrouped the output.

i Summaries were computed grouped by state and timePeriod.

i Output is grouped by state.

i Use `summarise(.groups = "drop\_last")` to silence this message.

i Use `summarise(.by = c(state, timePeriod))` for per-operation grouping
(`?dplyr::dplyr\_by`) instead.

A tibble: 107 x 3

Groups: state [54]

	state	timePeriod	totProtesters
	<chr>	<chr>	<int>
1	AK	2017_2019	30781
2	AK	2020	5196
3	AL	2017_2019	26629
4	AL	2020	7710
5	AR	2017_2019	15652
6	AR	2020	6197
7	AZ	2017_2019	189960
8	AZ	2020	33223
9	CA	2017_2019	2797687
10	CA	2020	391038

i 97 more rows

The result is still “grouped” at this point. Be sure to `ungroup()` so that subsequent calculations do not just occur within each group. The sequence `group()/summarize()/ungroup()`

is so common that there is a shortcut using `.by`. Either `ungroup()` or use `.by`. Then we `pivot_wider()`.

```
dataProtest |>
  filter(Date < "2021-01-01") |>
  mutate(year = substring(Date, 1, 4) |> as.numeric(),
         timePeriod = ifelse(year %in% 2017:2019,
                              "2017_2019",
                              "2020")) |>
  summarize(totProtesters = sum(Attendees, na.rm=TRUE),
            .by = c(state, timePeriod)) |>
  pivot_wider(names_from = timePeriod,
              names_prefix = "year",
              values_from = totProtesters,
              values_fill = 0) |>
  arrange(state)
```

```
# A tibble: 54 x 3
  state year2017_2019 year2020
  <chr>         <int>    <int>
1 AK             30781      5196
2 AL             26629      7710
3 AR             15652      6197
4 AZ             189960     33223
5 CA            2797687    391038
6 CO             382137     44987
7 CT             75083     30427
8 DC            1393990     64540
9 DE              6935      4195
10 FL            331471     81351
# i 44 more rows
```

Finally, we are in a position to compute the percent change, sort, and display those states with the largest percentage change.

```
dataProtest |>
  filter(Date < "2021-01-01") |>
  mutate(year = substring(Date, 1, 4) |> as.numeric(),
         timePeriod = ifelse(year %in% 2017:2019,
                              "2017_2019",
                              "2020")) |>
  summarize(totProtesters = sum(Attendees, na.rm=TRUE),
```

```

        .by = c(state, timePeriod)) |>
pivot_wider(names_from = timePeriod,
            names_prefix = "year",
            values_from = totProtesters,
            values_fill = 0) |>
mutate(ave2017_2019 = year2017_2019/3,
       pctChange = 100*(year2020-ave2017_2019)/ave2017_2019) |>
select(state, pctChange) |>
slice_max(pctChange, n = 5)

```

```

# A tibble: 5 x 2
  state pctChange
  <chr>     <dbl>
1 ND      180.08
2 SC      153.41
3 VA      128.58
4 LA      126.00
5 MS      125.13

```

15 Graphics and plots

We will finish our introduction to R by exploring **Tags** a little more through some barplots and a word cloud.

I will start with a special version of `mostCommon()` that will take a collection of tags and return the most common tag. This will allow us to find the most common protest type in the dataset. This function splits up the tags as we did before, and then applies `mostCommon()` to the resulting collection of tags.

```

mostCommonType <- function(x)
{
  reasons <- strsplit(x, "; ")
  reasons <- unlist(reasons)

  return( mostCommon(reasons) )
}

# test it out
dataProtest$Tags[1:10]

```

```
[1] "Healthcare; For Affordable Care Act"
[2] "Civil Rights; For racial justice; Martin Luther King, Jr."
[3] "Environment; For wilderness preservation"
[4] "Civil Rights; For racial justice; Martin Luther King, Jr."
[5] "Healthcare; For Planned Parenthood"
[6] "Executive; Against 45th president"
[7] "Executive; For 45th president; Counter protest"
[8] "Civil Rights; For racial justice; Against invited speaker"
[9] "Executive; Against 45th president"
[10] "Civil Rights; For women's rights; Women's March"
```

```
mostCommonType(dataProtest$Tags[1:10])
```

```
[1] "Civil Rights"
```

Now we can use `mostCommonType()` to find the most common protest type in the dataset. Note that `mostCommonType()` can return more than one value. `summarize()` will complain if it gets more than one value.

```
dataProtest |>
  group_by(state) |>
  summarize(mostCommonType(Tags)) |>
  print(n=Inf)
```

```
Error in `summarize()`:
i In argument: `mostCommonType(Tags)`.
i In group 12: `state = "GU"`.
Caused by error:
! `mostCommonType(Tags)` must be size 1, not 2.
i To return more or less than 1 row per group, use `reframe()``.
```

Why is it complaining about Guam? `separate_longer_delim()` can help here.

```
dataProtest |>
  filter(state == "GU") |>
  separate_longer_delim(Tags, delim = "; ") |> # from tidyr
  count(Tags) |>
  arrange(desc(n))
```

	Tags	n
1	Civil Rights	6
2	Other	6
3	Against war	3
4	Against abortion rights	2
5	Against development	2
6	Education	2
7	For historic preservation	2
8	Immigration	2
9	International	2
10	Military	2
11	Pride	2
12	Pro-LGBTQ	2
13	Against 45th president	1
14	Against colonization	1
15	Against hazardous conditions	1
16	Against increased prices	1
17	Against pandemic intervention	1
18	Against sexual/domestic violence	1
19	Coronavirus	1
20	Environment	1
21	Executive	1
22	For abortion rights	1
23	For compassionate immigration	1
24	For criminal justice	1
25	For greater accountability	1
26	For greater gun control	1
27	For refugees	1
28	For wilderness preservation	1
29	Guns	1
30	Healthcare	1
31	North Korea	1
32	Police	1
33	Tuition	1

Civil Rights and Other are themes in six protests... a tie.

So let's redo that with `reframe()` instead. `reframe()` is like `summarize()` but allows for multiple values.

```
dataProtest |>
  group_by(state) |>
```



```
reframe(mostCommonType(Tags)) |>
print(n=Inf)
```

```
# A tibble: 58 x 2
  state `mostCommonType(Tags)`
  <chr> <chr>
1 AK    Civil Rights
2 AL    Civil Rights
3 AR    Civil Rights
4 AZ    Civil Rights
5 CA    Civil Rights
6 CO    Civil Rights
7 CT    Civil Rights
8 DC    Civil Rights
9 DE    Civil Rights
10 FL   Civil Rights
11 GA    Civil Rights
12 GU    Civil Rights
13 GU    Other
14 HI    Other
15 IA    Civil Rights
16 ID    Civil Rights
17 IL    Civil Rights
18 IN    Civil Rights
19 KS    Civil Rights
20 KY    Civil Rights
21 LA    Civil Rights
22 MA    Civil Rights
23 MD    Civil Rights
24 ME    Civil Rights
25 MI    Civil Rights
26 MN    Civil Rights
27 MO    Civil Rights
28 MS    Civil Rights
29 MT    Civil Rights
30 NC    Civil Rights
31 ND    Civil Rights
32 NE    Civil Rights
33 NH    Civil Rights
34 NJ    Civil Rights
35 NM    Civil Rights
36 NV    Civil Rights
```

```

37 NY    Civil Rights
38 OH    Civil Rights
39 OK    Civil Rights
40 OR    Civil Rights
41 PA    Civil Rights
42 PR    Against corruption
43 PR    Against state executive
44 PR    Executive
45 RI    Civil Rights
46 SC    Civil Rights
47 SD    Civil Rights
48 Space Against 45th president
49 Space Executive
50 TN    Civil Rights
51 TX    Civil Rights
52 UT    Civil Rights
53 VA    Civil Rights
54 VT    Civil Rights
55 WA    Civil Rights
56 WI    Civil Rights
57 WV    Civil Rights
58 WY    Civil Rights

```

So why does Puerto Rico show up three times in these results?

```

dataProtest |>
  filter(state == "PR") |>
  separate_longer_delim(Tags, delim = "; ") |> # from tidyr
  count(Tags) |>
  slice_max(n)

```

```


      Tags  n
1  Against corruption 11
2 Against state executive 11
3      Executive 11

```

There are three tags all with 11 protests each, a three-way tie for the largest number of protests. So `mostCommonType()` returns all three tags.


R has a lot of built-in functions for creating plots and graphics. We will use the `barplot()` function to create a bar plot of the average number of attendees at protests in each state.

```
a <- dataProtest |>
  group_by(state) |>
  summarize(Attendees=mean(Attendees, na.rm=TRUE))
barplot(a$Attendees, names.arg = a$state)
```




The state name labels are too big so we can shrink the “character expansion” (`cex`) by half.

```
barplot(a$Attendees, names.arg = a$state, cex.names=0.5)
```




We can also make the plot horizontal.

```
barplot(a$Attendees, names.arg = a$state,  
       cex.names=0.3,  
       horiz=TRUE,  
       col="seagreen",  
       xlim=c(0,5000))
```



We can also create a bar plot of the number of protests for the top 5 reasons.

```
reasons <- dataProtest$Tags |>
  strsplit("; ") |>
  unlist() |>
  table() |>
  sort(decreasing = TRUE) |>
  head(5)
barplot(reasons,
  ylab="Number of Protests",
  xlab="Protest Reason",
  cex.names = 0.6) # shrink bar labels 30%
```




For figures and plots, always use a vector graphics format. That means export your graphics using SVG or EMF. These formats are scalable and will look good at any size. You can insert these graphics into Word, PowerPoint, or Google Docs. PNG graphics tend to look blurry in reports and presentations. Show some pride in your data work by making sure that your final product looks great. Stick with SVG or EMF or another vector graphics format.

We will end with a beautiful word cloud of the protest tags.

```
library(wordcloud2)
```

Warning: package 'wordcloud2' was built under R version 4.6.1

```
dataProtest$Tags |>  
  strsplit(split="; ") |>  
  unlist() |>  
  table() |>  
  wordcloud2()
```



3

16 Review

As you saw in this script, R has a lot of functions. We started off figuring out how to set our file path so R knows where to look for files. We loaded the data from a .RData file and we listed all the objects in R's environment.

- `setwd()` set working directory
- `load()` load R objects saved in a .RData file
- `ls()` list objects in the R environment

R, of course, has all the basic math operations that you might need to do with a set of numbers. Like

- `sqrt()`
- `log()`, note that `log()` is the natural log as it is in most mathematical programming languages
- `round()` round to the nearest integer
- `abs()` absolute value
- `length()` number of elements in a collection
- `cumsum()` cumulative sum
- `sum()`, `mean()`, `median()`, `min()`, `max()`

Then we worked through some basic functions to work with R objects.

- `c()` combine numbers and other R objects together in a collection
- `[]` index an object to pick out selected values, rows, or columns
- `nrow()`, `ncol()`
- `head()`, `tail()`

When working with datasets, we covered all the standard functions needed to manipulate data.

- `slice()`, `slice_max()`, `slice_min()` pick out rows by their position in the dataset or by the max/min values
- `filter()` pick out rows based on a logical expression about what is in that row
- `select()` pick out columns by name
- `count()` count the number of rows in a dataset or the number of rows in a dataset by groups
- `mutate()` create new columns or edit existing columns
- `str_sub()` extract substrings from a string
- `replace_values()` used inside `mutate()` to change selected values in a column
- `group_by()`, `summarize()`, `reframe()` used to summarize data by groups
- `arrange()` sort rows in a dataset
- `pivot_wider()` and `pivot_longer()` to organize data in wide or long form

We also covered some more advanced functions.

- `grep1()` search for patterns in text
- `summary()` get a summary of a dataset or any set of numbers
- `sort()` sort a collection of numbers
- `unlist()` remove the list structure from a list
- `names()` get the names of the elements in a collection
- `as.numeric()` convert objects to numbers, we specifically converted logical values to 1s and 0s
- `strsplit()` split a string into a list of substrings

And we made some graphics too.

- `barplot()` create a bar plot
- `wordcloud2()` create a word cloud

In addition we even created our own new functions!

- `mostCommon()` find the most common value in a collection
- `mostCommonType()` find the most common tag in a string containing semi-colon separated tags

Before looking at the solutions, try out the exercises for yourself. All the skills you will be learning build on the fundamentals presented in this script. It would be a good idea to go through this a second time to make sure you understand everything.

17 Solutions to the exercises

1. What is the date of the protest in line 10000 of the dataset?

```
dataProtest |>
  slice(10000) |>
  select(Date)
```

```
      Date
1 2018-03-24
```

2. Which protest type is in line 4289 of the dataset?

```
dataProtest |>
  slice(4289) |>
  select(Tags)
```

Tags

1 International; For Palestine; Israel

3. How many protests occurred in your home state?

```
dataProtest |>
  filter(state == "CA") |>
  count()
```

```
      n
1 4439
```

4. Where did the protest in the last row of the full dataset occur?

```
dataProtest |>
  select(state, Location) |>
  tail(1)
```

```
      state      Location
38096    CA San Francisco, CA
```

5. Which state had the most protests?

```
dataProtest |>
  count(state) |>
  slice_max(n,
            with_ties = TRUE) # in case of ties
```

```
state      n
1    CA 4439
```

6. Which state had the least protests?

```
dataProtest |>
  count(state) |>
  slice_min(n, with_ties = TRUE)
```

```
state n
1 Space 1
```

7. Which state had the most civil rights protests?

```
dataProtest |>
  filter(civilrights==1) |>
  count(state) |>
  slice_max(n, with_ties = TRUE)
```

```
state    n
1      CA 1424
```

8. Create a new column that is 1 if the protest has the tag 'Against pandemic intervention'

```
dataProtest <- dataProtest |>
  mutate(pandemic = as.numeric(grepl("Against pandemic intervention", Tags)))
```

9. Which state had the most protests against pandemic interventions?

```
dataProtest |>
  filter(pandemic == 1) |>
  count(state) |>
  slice_max(n, with_ties = TRUE)
```

```
state    n
1      CA 227
```

10. Are civil rights protests larger on average than non-civil rights protests?

```
dataProtest |>
  group_by(civilrights) |>
  summarize(mean(Attendees, na.rm=TRUE))
```

```
# A tibble: 2 x 2
  civilrights `mean(Attendees, na.rm = TRUE)`
    <dbl>          <dbl>
1       0       342.17
2       1      1113.0
```

```
# Yes, civil rights protests are larger on average than non-civil rights protests.
```